

Alex Look

DSC 324

Homework Module 2

1/23/22

1a) Copy of dataset with select columns

```
mtcars.cyl      mtcars.disp      mtcars.hp
Min.   :4.000    Min.   : 71.1    Min.   : 52.0
1st Qu.:4.000    1st Qu.:120.8    1st Qu.: 96.5
Median :6.000    Median :196.3    Median :123.0
Mean   :6.188    Mean   :230.7    Mean   :146.7
3rd Qu.:8.000    3rd Qu.:326.0    3rd Qu.:180.0
Max.   :8.000    Max.   :472.0    Max.   :335.0

mtcars.wt      mtcars.carb
Min.   :1.513    Min.   :1.000
1st Qu.:2.581    1st Qu.:2.000
Median :3.325    Median :2.000
Mean   :3.217    Mean   :2.812
3rd Qu.:3.610    3rd Qu.:4.000
Max.   :5.424    Max.   :8.000
```

1b) Column of 1's named "count"

Description: df [32 x 6]

count <dbl>	mtcars.cyl <dbl>	mtcars.disp <dbl>	mtcars.hp <dbl>	mtcars.wt <dbl>	mtcars.carb <dbl>
1	6	160.0	110	2.620	4
1	6	160.0	110	2.875	4
1	4	108.0	93	2.320	1
1	6	258.0	110	3.215	1
1	8	360.0	175	3.440	2
1	6	225.0	105	3.460	1
1	8	360.0	245	3.570	4
1	4	146.7	62	3.190	2
1	4	140.8	95	3.150	2
1	6	167.6	123	3.440	4
1	6	167.6	123	3.440	4
1	8	275.8	180	4.070	3
1	8	275.8	180	3.730	3

1-13 of 32 rows

Previous **1** 2 3 Next

1c) Convert dataset into matrix

	count	mtcars.cyl	mtcars.disp	mtcars.hp	mtcars.wt
[1,]	1	6	160.0	110	2.620
[2,]	1	6	160.0	110	2.875
[3,]	1	4	108.0	93	2.320
[4,]	1	6	258.0	110	3.215
[5,]	1	8	360.0	175	3.440
[6,]	1	6	225.0	105	3.460
[7,]	1	8	360.0	245	3.570
[8,]	1	4	146.7	62	3.190
[9,]	1	4	140.8	95	3.150
[10,]	1	6	167.6	123	3.440
[11,]	1	6	167.6	123	3.440
[12,]	1	8	275.8	180	4.070
[13,]	1	8	275.8	180	3.730
[14,]	1	8	275.8	180	3.780
[15,]	1	8	472.0	205	5.250
[16,]	1	8	460.0	215	5.424
[17,]	1	8	440.0	230	5.345
[18,]	1	4	78.7	66	2.200
[19,]	1	4	75.7	52	1.615
[20,]	1	4	71.1	65	1.835
[21,]	1	4	120.1	97	2.465
[22,]	1	8	318.0	150	3.520
[23,]	1	8	304.0	150	3.435
[24,]	1	8	350.0	245	3.840
[25,]	1	8	400.0	175	3.845
[26,]	1	4	79.0	66	1.935
[27,]	1	4	120.3	91	2.140
[28,]	1	4	95.1	113	1.513
[29,]	1	8	351.0	264	3.170
[30,]	1	6	145.0	175	2.770
[31,]	1	8	301.0	335	3.570
[32,]	1	4	121.0	109	2.780

	mtcars.carb
[1,]	4
[2,]	4
[3,]	1
[4,]	1
[5,]	2
[6,]	1
[7,]	4
[8,]	2
[9,]	2
[10,]	4
[11,]	4
[12,]	3
[13,]	3
[14,]	3
[15,]	4
[16,]	4
[17,]	4
[18,]	1
[19,]	2
[20,]	1
[21,]	1
[22,]	2
[23,]	2
[24,]	4
[25,]	2
[26,]	1
[27,]	2
[28,]	2
[29,]	4
[30,]	6
[31,]	8
[32,]	2

	mtcars.mpg
[1,]	21.0
[2,]	21.0
[3,]	22.8
[4,]	21.4
[5,]	18.7
[6,]	18.1
[7,]	14.3
[8,]	24.4
[9,]	22.8
[10,]	19.2
[11,]	17.8
[12,]	16.4
[13,]	17.3
[14,]	15.2
[15,]	10.4
[16,]	10.4
[17,]	14.7
[18,]	32.4
[19,]	30.4
[20,]	33.9
[21,]	21.5
[22,]	15.5
[23,]	15.2
[24,]	13.3
[25,]	19.2
[26,]	27.3
[27,]	26.0
[28,]	30.4
[29,]	15.8
[30,]	19.7
[31,]	15.0
[32,]	21.4

1d) Compute beta coefficients

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
count	1.00	1.000	1.00	1.000	1.00	1.00
mtcars.cyl	6.00	6.000	4.00	6.000	8.00	6.00
mtcars.disp	160.00	160.000	108.00	258.000	360.00	225.00
mtcars.hp	110.00	110.000	93.00	110.000	175.00	105.00
mtcars.wt	2.62	2.875	2.32	3.215	3.44	3.46
mtcars.carb	4.00	4.000	1.00	1.000	2.00	1.00
	[,7]	[,8]	[,9]	[,10]	[,11]	[,12]
count	1.00	1.00	1.00	1.00	1.00	1.00
mtcars.cyl	8.00	4.00	4.00	6.00	6.00	8.00
mtcars.disp	360.00	146.70	140.80	167.60	167.60	275.80
mtcars.hp	245.00	62.00	95.00	123.00	123.00	180.00
mtcars.wt	3.57	3.19	3.15	3.44	3.44	4.07
mtcars.carb	4.00	2.00	2.00	4.00	4.00	3.00
	[,13]	[,14]	[,15]	[,16]	[,17]	[,18]
count	1.00	1.00	1.00	1.000	1.000	1.0
mtcars.cyl	8.00	8.00	8.00	8.000	8.000	4.0
mtcars.disp	275.80	275.80	472.00	460.000	440.000	78.7
mtcars.hp	180.00	180.00	205.00	215.000	230.000	66.0
mtcars.wt	3.73	3.78	5.25	5.424	5.345	2.2
mtcars.carb	3.00	3.00	4.00	4.000	4.000	1.0
	[,19]	[,20]	[,21]	[,22]	[,23]	[,24]
count	1.000	1.000	1.000	1.00	1.000	1.00
mtcars.cyl	4.000	4.000	4.000	8.00	8.000	8.00
mtcars.disp	75.700	71.100	120.100	318.00	304.000	350.00
mtcars.hp	52.000	65.000	97.000	150.00	150.000	245.00
mtcars.wt	1.615	1.835	2.465	3.52	3.435	3.84
mtcars.carb	2.000	1.000	1.000	2.00	2.000	4.00
	[,25]	[,26]	[,27]	[,28]	[,29]	[,30]
count	1.000	1.000	1.00	1.000	1.00	1.00
mtcars.cyl	8.000	4.000	4.00	4.000	8.00	6.00
mtcars.disp	400.000	79.000	120.30	95.100	351.00	145.00
mtcars.hp	175.000	66.000	91.00	113.000	264.00	175.00
mtcars.wt	3.845	1.935	2.14	1.513	3.17	2.77
mtcars.carb	2.000	1.000	2.00	2.000	4.00	6.00
	[,31]	[,32]				
count	1.00	1.00				
mtcars.cyl	8.00	4.00				
mtcars.disp	301.00	121.00				
mtcars.hp	335.00	109.00				
mtcars.wt	3.57	2.78				
mtcars.carb	8.00	2.00				

1e) Comparing manual calculation to output of function

```
call:
lm(formula = mpg ~ cyl + disp + hp + wt + carb, data = mtcars)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-4.0635 -1.4580 -0.4306  1.2927  5.8244
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  40.815359   3.025568  13.490   3e-13 ***
cyl          -1.291899   0.679227  -1.902   0.06830 .
disp           0.011486   0.015375   0.747   0.46175
hp            -0.020353   0.020062  -1.015   0.31968
wt            -3.846949   1.192155  -3.227   0.00337 **
carb          -0.006747   0.574269  -0.012   0.99072
```

Signif. codes:

0 '***', 0.001 '**', 0.01 '*', 0.05 '.', 0.1 ' ', 1

Residual standard error: 2.56 on 26 degrees of freedom
Multiple R-squared: 0.8486, Adjusted R-squared: 0.8195
F-statistic: 29.15 on 5 and 26 DF, p-value: 7.056e-10

```
      mtcars.mpg
[1,] 40.815359236
[2,] -1.291898563
[3,]  0.011485584
[4,] -0.020352893
[5,] -3.846949031
[6,] -0.006746893
```

When looking at the manual calculations I did compared to the ones that was outputted to the function, the results that were returned were exact same numbers in which confirms that the result I hand calculated was the same as what the function had outputted. I will not consider rounding because that is something very meniscal to the idea of what the problem had asked because the results are vastly similar.

2a) Running regression model

Training MEDV Adj R²: .7491.

Root MSE	4.66119	R-Square	0.7571
Dependent Mean	22.30737	Adj R-Sq	0.7491
Coeff Var	20.89531		

RSME of both Training and Test sets:

Residual standard error: 4.661 on 367 degrees of freedom
Multiple R-squared: 0.7571, Adjusted R-squared: 0.7491
F-statistic: 95.31 on 12 and 367 DF, p-value: < 2.2e-16


```
[1] 4.580769  
[1] 5.263608
```

Housing Training RMSE: 4.580769

Housing Testing RMSE: 5.263608

I believe that there is no evidence of overfitting. Between both training and testing RMSE's there is very little difference between the two values.

2b)



When looking at the graph there is no obvious dip in this graph so finding a minimum lambda might not exist.


2c)

```
[1] 63.53669
```

The dev value used to be 75 and now is 20.61% that would mean I am losing coefficients and the R-value was lost so it may not have regularized correctly.

2d)

	Df	%Dev	Lambda
1	12	20.61	69.73



The Lasso is doing better cause of the dev percentage difference is less than the ridge.

2e)

	Df	%Dev	Lambda
1	1	19.48	3.809

There is one variable selected in the lambda.1se under Df. The variance after the OLS went down even more than the Lasso or Ridge to 19.48 from the initial 75%.

2f) OLS, Ridge, Lasso

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.965e+01	5.610e+00	7.069	7.94e-12	***
CRIM	-1.299e-01	3.412e-02	-3.807	0.000165	***
ZN	4.341e-02	1.570e-02	2.764	0.005994	**
INDUS	6.302e-03	6.958e-02	0.091	0.927884	
CHAS	3.594e+00	9.454e-01	3.802	0.000168	***
NOX	-2.197e+01	4.377e+00	-5.021	8.05e-07	***
RM	4.229e+00	4.898e-01	8.634	< 2e-16	***
AGE	-1.268e-04	1.511e-02	-0.008	0.993307	
DIS	-1.529e+00	2.318e-01	-6.598	1.46e-10	***
RAD	2.665e-01	7.341e-02	3.630	0.000324	***
TAX	-1.134e-02	4.130e-03	-2.746	0.006338	**
PTRATIO	-9.828e-01	1.506e-01	-6.526	2.24e-10	***
LSTAT	-4.665e-01	6.094e-02	-7.655	1.73e-13	***

Signif. codes:

0 '***', 0.001 '**', 0.01 '*', 0.05 '.', 0.1 ' ' 1

Residual standard error: 4.661 on 367 degrees of freedom

Multiple R-squared: 0.7571, Adjusted R-squared: 0.7491

F-statistic: 95.31 on 12 and 367 DF, p-value: < 2.2e-16

13 x 1 sparse Matrix of class "dgCMatrix"

```
      s1
(Intercept) -0.122160867
ZN           -0.002991104
INDUS        0.033262853
CHAS        -0.195382591
NOX          2.282376267
RM          -0.120336124
AGE          0.006961588
DIS         -0.121059676
RAD          0.056792060
TAX          0.002498431
PTRATIO      0.083324303
LSTAT        0.036690422
MEDV        -0.027024183
```

13 x 1 sparse Matrix of class "dgCMatrix"

```
      s1
(Intercept) 1.8617434
ZN           .
INDUS        .
CHAS         .
NOX          .
RM           .
AGE          .
DIS          .
RAD          0.1964528
TAX          .
PTRATIO      .
LSTAT        .
MEDV         .
```

When looking at all the different datasets OLS and Ridge are considerably close in coefficients. Lasso only has 1 coefficient and that value compared to the other two are not very close.

2g) If I had to pick a model to use out of these, I would use the OLS model because the values would already be checked for overfitting. Each model after Ridge and Lasso we are losing coefficients after each run and to me I would rather have the whole dataset and it already be checked for overfitting.

3a) Predicting the test set

Residual standard error: 1.303 on 17 degrees of freedom
Multiple R-squared: 0.9272, Adjusted R-squared: 0.9058
F-statistic: 43.32 on 5 and 17 DF, p-value: 4.387e-09


[1] 1.120282

[1] 4.260169

I feel there is overfitting because of the huge gap in RMSE from the training and test sets.

3b)


[1] 15.37505



That 15.37505 is the RMSE.

3c)


[1] 15.37505




It is doing better because there is less lambda but the RMSE is the same value between both.

3d)

3e) The .5 RMSE was the closest RMSE to the test RMSE. All three were better than the ridge and lasso.





3f) When mixing both ridge and lasso its better than the extremes to have regularized seeing as though .5 was much closer to the test RMSE.